

Numerical Python for Scalable Architectures

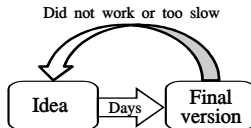
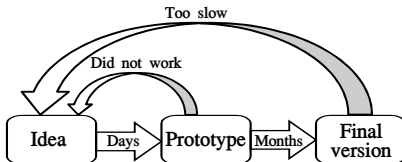
Mads Ruben Burgdorff Kristensen
Brian Vinter

eScience Centre
University of Copenhagen
Denmark
madsbk@diku.dk/vinter@diku.dk

October 14, 2010

- Motivation
- Python
- Numerical Python (NumPy)
- DistNumPy
- Benchmarks

Motivation – Workflow



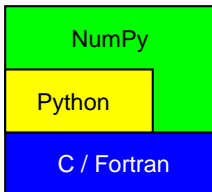
- High Productivity
 - High-level language
 - No compilation
 - Interactive
- High Performance
 - Low-level language such as C, Fortran, etc.
 - Compiling to machine code
 - Parallel Programming

Python

- General-purpose high-level programming language
- The design philosophy emphasizes code readability.
- Sacrificing performance over productivity
- Gluing libraries together

Numerical Python (NumPy)

- Framework for numerical computation similar to Matlab
- Introduces efficient arrays and a lot of useful array operations
 - Linear algebra functions
 - Fourier transforms
 - Random number arrays



Vectorized operations called Universal Functions

```
#C
for(i = 0; i < rows; i++)
    c[i] = a[i] + b[i];
```

```
#NumPy
c = a + b
```

```
#Python
c = []
for i in range(len(a)):
    c.append(a[i] + b[i])
```

Matching arrays that have different dimensions

```
#No broadcast
for i in range(len(A)):
    c[i] = a[i] + b
```

```
#Broadcast
c = a + b
```

05	06		10	20		15	26
03	04	+	10	20	=	13	04
01	02		10	20		11	22

--

 = Broadcasted element

DistNumPy

A distributed version of NumPy

Ideal workflow – High Productivity and High Performance

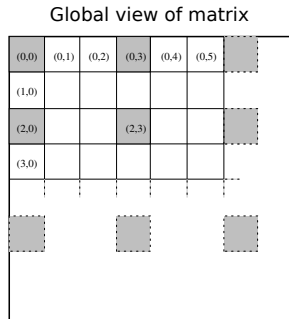
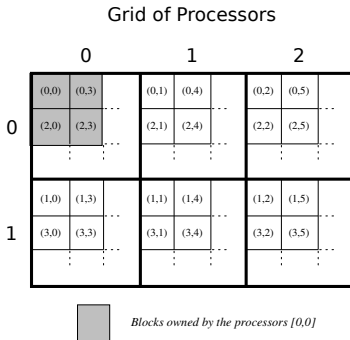
Transparent parallelism

- DistNumPy introduces a distributed parallel array-backend
- Still sequential programming
- Parallel execution of universal functions

Monte Carlo π simulation

```
from numpy import *
S = 1000 #Number of samples
(x, y) = (empty([S], dist=True), empty([S], dist=True))
(x, y) = (random(x), random(y))
(x, y) = (square(x), square(y))
z = (x + y) < 1
print add.reduce(z) * 4.0 / S #The result
```

N-Dimensional Block Cyclic Distribution



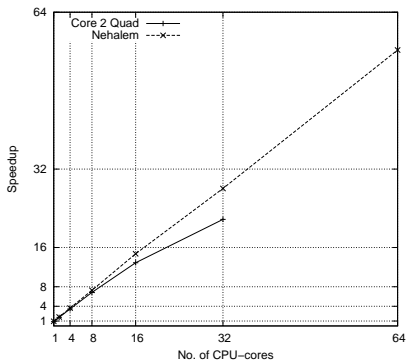
- Used in High Performance Fortran
- Diagonal workflow e.g. Gaussian elimination

- DistNumPy make use of MPI version 2.1
- One-sided, two-sided and collective communication
- Latency hiding by double buffering

Table: Hardware specifications

CPU	Core 2 Quad	Nehalem
CPU Frequency	2.26 GHz	2.66 GHz
CPU per node	1	2
Cores per CPU	4	4
Memory per node	8 GB @ 6.5 GB/s	24 GB @ 25.6 GB/s
Number of nodes	8	8
Network	Gigabit Ethernet	Gigabit Ethernet

Monte Carlo π simulation

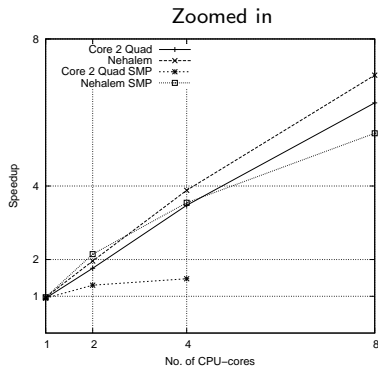
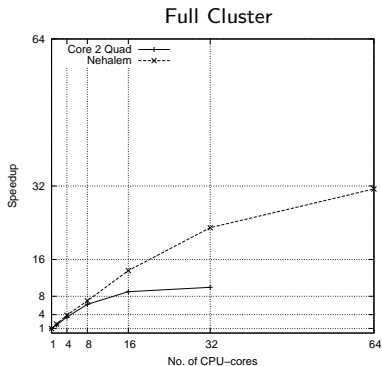


Nehalem – CPU utilization of 88% on 64 CPU-cores

Jacobi.py

```
h = zeros(shape(B), float, dist=True)
dmax = 1.0
AD = A.diagonal()
while(dmax > tol):
    hnew = h + (B - add.reduce(A * h, 1)) / AD
    tmp = absolute((h - hnew) / h)
    dmax = maximum.reduce(tmp)
    h = hnew
print h #The result
```

Scalability – Jacobi Solver

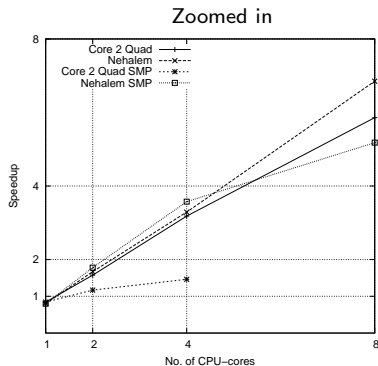
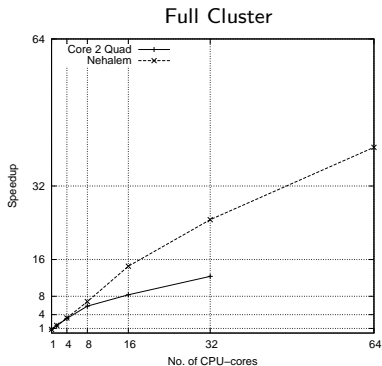


Nehalem – CPU utilization of 85% on 16 CPU-cores and 50% on 64 CPU-cores

N-body simulation

```
for i in range(k):  
    Fx = dot(OnesCol, PxT) - dot(Px, OnesRow)  
    Dsq = Fx * Fx + Fy * Fy + Fz * Fz + Identity  
    D = sqrt(Dsq)  
    #mutual forces between all pairs of objects  
    F = G * dot(M, MT) / Dsq  
    F = F - diag(diag(F))#set 'self attraction' to 0  
    Fx = (Fx / D) * F  
    #net force on each body  
    Fnet_x = add.reduce(Fx,1)  
    Fnet_x = Fnet_x[:,newaxis]  
    Fnet_x *= dT  
    #change in velocity:  
    Vx += Fnet_x / M  
    #change in position  
    Px += Vx * dT
```


Scalability – N-body simulation



Nehalem – CPU utilization of 91% on 16 CPU-cores and 63% on 64 CPU-cores

- Fully transparent data distribution
- Fully transparent parallel execution
- However, the use of Universal Functions is required
- DistNumPy running Jacobi is roughly 50% slower than the C
 - 21 seconds for C
 - 31 seconds for NumPy
 - 32 seconds for DistNumPy (17 seconds on two CPU-cores)